


Two independent lines of evidence for vision-token redundancy

Uniform (position-only) token pruning barely hurts task performance



Vision tokens are low-frequency redundant in two unrelated encoders

